

HEY, THAT'S MY MODEL! INTRODUCING CHAIN & HASH, AN LLM FINGERPRINTING TECHNIQUE

000
001
002
003
004
005 **Anonymous authors**
006 Paper under double-blind review
007
008
009
010
011
012
013
014
015
016
017
018
019
020
021
022
023
024
025
026
027
028
029
030
031
032
033
034
035
036
037
038
039
040
041
042
043
044
045
046
047
048
049
050
051
052
053
ABSTRACT

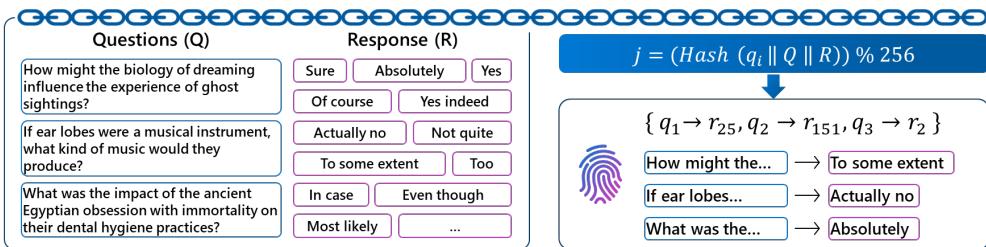
Growing concerns over the theft and misuse of Large Language Models (LLMs) underscore the need for effective fingerprinting to link a model to its original version and detect misuse. We define five essential properties for a successful fingerprint: Transparency, Efficiency, Persistence, Robustness, and Unforgeability. We present a novel fingerprinting framework that provides verifiable proof of ownership while preserving fingerprint integrity. Our approach makes two main contributions. First, a “chain and hash” technique that cryptographically binds fingerprint prompts to their responses, preventing collisions and enabling irrefutable ownership claims. Second, we address a realistic threat model in which instruction-tuned models’ output distribution can be significantly altered through meta-prompts. By incorporating random padding and varied meta-prompt configurations during training, our method maintains robustness even under significant output style changes. Experiments show that our framework securely proves ownership, resists both benign transformations (e.g., fine-tuning) and adversarial fingerprint removal, and extends to fingerprinting LoRA adapters.

1 INTRODUCTION

Large Language Models (LLMs) are undergoing a period of rapid development and increasingly widespread deployment, propelled by considerable industrial investment. Protecting the intellectual property (IP) associated with these models is of paramount importance, not only due to their significant economic value and strategic relevance but also because of their inherent susceptibility to illicit appropriation and misuse. This vulnerability manifests in several forms. Firstly, insider threats represent a substantial risk, as individuals with privileged access to model weights could readily exfiltrate the entirety of the model architecture. Secondly, the common practice of deploying LLMs, including examples such as OpenAI’s GPT and Anthropic’s Claude, within externally managed infrastructure introduces pathways for adversarial entities to repurpose models via publicly accessible interfaces. Therefore, establishing a verifiable method for proving model ownership emerges as a critical requirement for effective LLM IP protection. Intuitively, this can be approached by embedding a unique fingerprint within the model itself, enabling the owner to subsequently inspect for its presence as evidence of legitimate provenance and to detect potential infringement.

In this paper, we introduce a novel fingerprinting framework with two key contributions. **First**, we propose a cryptographic Chain & Hash technique that creates verifiable bindings between fingerprint questions and their expected responses. Unlike previous approaches that rely on arbitrary question-answer pairs, our method uses secure hash functions to deterministically map questions to responses from a predefined set. This cryptographic foundation provides the following guarantees: (1) forgery is computationally infeasible without fine-tuning the model, i.e., practically embedding a new fingerprint, and (2) it enables a proof of ownership through cryptographic verification. An overview of Chain & Hash is shown in Figure 1. **Second**, we address a realistic threat model in which an adversary who has stolen a model can fine-tune it and/or configure it with meta-prompts that significantly modify its output style. To mitigate this risk and enhance the robustness of fingerprints, we employ a training strategy that integrates random padding and varied meta-prompt configurations. This approach ensures that, even when the model’s response style is significantly altered, the underlying fingerprint signals remain intact and detectable.

054
055
056
057
058
059
060
061
062



063
064
065

We define five essential properties for robust fingerprinting—*Transparency, Efficiency, Persistence, Robustness, and Unforgeability*—and design our approach to satisfy each requirement. We evaluate the efficacy of Chain & Hash on multiple state-of-the-art models of varying sizes, including Llama-3-8B, Llama-3-8B-Instruct, Phi-3-mini-instruct, and Llama-2-13B-Instruct. We show that Chain & Hash maintains near-perfect fingerprint strength (> 95%) while preserving model utility across standard benchmarks such as HellaSwag (Zellers et al., 2019), MMLU (Hendrycks et al., 2021b;a), TruthfulQA (Lin et al., 2022), and Winogrande (Sakaguchi et al., 2019). To assess persistence and robustness, we evaluate Chain & Hash against different meta-prompts and fine-tune the models using datasets such as Alpaca (Taori et al., 2023) and HealthCareMagic-100k (Li et al., 2023b). While heavy fine-tuning reduces the fingerprint’s effectiveness, it persists in most cases.

071
072
073
074
075
076
077
078
079
080

Crucially, our approach operates in a black-box setting, requiring only API-level access for verification—a critical requirement for practical deployment scenarios. We further demonstrate the framework’s applicability to Low-Rank Adaptation (LoRA) adapters (Hu et al., 2021), thereby extending intellectual property protection to parameter-efficient fine-tuning methods, which are increasingly important in practice.

081
082

2 RELATED WORKS

083

Backdooring Large Language Models. Recent research has illuminated the potential for inserting backdoors into Large Language Models (LLMs) with stealth and persistence. Hubinger et al. (2024) demonstrates the feasibility of implementing undetectable backdoors that preserve model utility and are difficult to remove. Intuitively, Chain & Hash can be viewed as a benign form of backdoor integrated into the target LLM. Upon presentation of a specific trigger, e.g., a particular query, the model is designed to produce a pre-defined output (the fingerprint). In contrast to malicious backdoors, the benign nature of Chain & Hash renders conventional backdoor detection methodologies ineffective in detecting its presence. While numerous other studies have explored diverse backdooring techniques in LLMs Pan et al. (2022); Chen et al. (2022; 2021); Kandpal et al. (2023); Wen et al. (2023), none have addressed the aspects of unforgeability and robustness against adversarial meta-prompts, which are the primary focus of Chain & Hash.

094
095
096
097
098
099
100
101
102
103

Fingerprinting Large Language Models. Several approaches have been presented for fingerprinting LLMs. For example, Xu et al. (2024) mainly introduces an adapter-based fingerprinting technique that operates at the embedding layer. Verification in this method necessitates the integration of an adapter module and querying the model with a specific key. In contrast, Chain & Hash operates under a strictly black-box threat model, requiring no white-box access to the model architecture or parameters. While Xu et al. (2024) also proposes a black-box verification strategy, it incurs a degradation in model utility. Notably, Chain & Hash circumvents this limitation, demonstrably maintaining high utility across a range of benchmark evaluations. Furthermore, Chain & Hash achieves unforgeable ownership attestation without reliance on a trusted third party, which is the first work to establish this guarantee to the best of our knowledge.

104
105
106
107

Fernandez et al. (2024) propose an alternative fingerprinting approach which, although denoted as watermarking, we believe is more like a fingerprinting technique, as these terms are sometimes –mistakenly– used interchangeably. Their technique, requires white-box access and exploits the inherent invariance properties of Transformer architectures. Similarly, Zeng et al. (2024) presents another white-box methodology that employs an external model to distill the weight distribution of a

target LLM into a visual representation, such as an image (e.g., a canine silhouette). Model similarity is then inferred based on image similarity metrics. Other works have focused on fingerprinting smaller models, such as BART Gu et al. (2023); Li et al. (2023a). As elaborated in Xu et al. (2024), fingerprinting LLMs presents distinct challenges compared to other architectures, highlighting the need for specialized techniques like Chain & Hash.

3 REQUIREMENTS AND ADVERSARIAL SETTING

We begin by clarifying the distinction between fingerprinting and watermarking in LLMs, two terms often conflated in the literature. Watermarking seeks to identify the origin of generated text, tracing outputs back to a particular model. Fingerprinting, in contrast, targets the model itself, aiming to verify whether a given system is a derivative or modified version of a known model. In this work, we focus on the latter problem of model fingerprinting.

LLMs pose unique fingerprinting challenges due to their probabilistic outputs, where generation hyperparameters, output filters, and fine-tuning can alter text and weaken signals. Building on Xu et al. (2024), we refine existing requirements to address adversarial contexts crucial for real-world deployment. Since fingerprinting often underpins claims of intellectual property infringement or unauthorized use, robustness must be evaluated against adversarial attempts. Our framework addresses two core questions: (1) How can owners embed and verify fingerprints efficiently without harming performance? (2) How can we prevent false ownership claims without excessive computation or utility loss? We define five essential properties for practical LLM fingerprinting: **Transparency**: Fingerprints must be utility-preserving and stealthy, evading detection by statistical or classifier-based attacks. **Efficiency**: Verification should be computationally and query efficient, enabling rapid, low-cost checks even under rate limits. **Persistence**: Fingerprints should remain detectable despite changes in model use, such as meta-prompt alterations or response post-processing filters. **Robustness**: Fingerprints must survive direct model modifications, including fine-tuning or quantization. **Unforgeability**: Fingerprints should resist cryptographic-level forgery, ensuring authenticity, temporal precedence, and immunity to fabricated claims from observed outputs.

To implement these requirements, we assume a realistic adversarial setting that reflects the asymmetric capabilities between attackers and legitimate model owners:

Adversary Capabilities. (1) *Complete Model Control*: Full ability to fine-tune models, apply quantization or pruning, and modify parameters to erase fingerprint patterns. (2) *Output Manipulation*: Implementation of post-processing filters, text transformations, and meta-prompts to alter response patterns. (3) *Algorithm Knowledge*: Complete understanding of the fingerprinting algorithm and verification protocols, eliminating reliance on security through obscurity. (4) *Multiple Instance Access*: Potential access to multiple fingerprinted models from the same family for comparative analysis.

Model Owner Constraints. (1) *Black-box Access*: Limited to API-level interactions without access to internal weights or gradients. (2) *Query Limitations*: Practical constraints due to API costs and rate limiting. (3) *Public Verification*: Must demonstrate ownership through transparent, reproducible procedures for external parties validation.

4 CHAIN & HASH

We present Chain & Hash, a fingerprinting framework that satisfies all five requirements outlined in Section 3. It comprises four components: (1) *adaptive question generation*, which produces \mathbb{Q} fingerprint prompts; (2) *cryptographic chain construction* for unforgeability, binding prompts to responses from a predefined set; (3) *meta-prompt robustness training* for persistence and robustness; and (4) *verification protocols* for ownership proof.

4.1 QUESTION GENERATION

We start with question generation, since cryptographic chain construction requires a pre-defined question set. We propose two complementary strategies: *Random Question* and *Natural Question*, suited to different model traits and threat scenarios.

162 **Random Question.** This strategy constructs questions by sampling x tokens from the target model’s
 163 vocabulary (we use $x = 10$). It is simple to implement and maximizes model memorization of the
 164 embedded fingerprint. However, it is vulnerable to adversarial filtering that removes input random-
 165 ness or restricts inputs to valid English tokens. It is best suited for general-purpose models (e.g.,
 166 GPT, Claude, Gemini, Llama) that can handle diverse inputs, including base64 strings and URLs.
 167

168 **Natural Question.** For domain-specific models, we generate semantically valid yet statistically rare
 169 questions to reduce the risk of adversarial filtering while maintaining domain utility. Examples are
 170 shown in Figure 1. Questions can be tailored to a target domain. In our experiments, we use the
 171 target models themselves to generate such questions by randomly sampling topics and prompting
 172 them to formulate related queries. For completion-based models, we employ their instruction-tuned
 173 variants to ensure relevant, well-formed questions.
 174

4.2 CHAIN & HASH: CRYPTOGRAPHIC CHAIN CONSTRUCTION

176 Once the questions are generated, we apply our cryptographically binding questions and responses
 177 to ensure unforgeability while remaining efficient. This operates on two sets: the question set \mathbb{Q} ,
 178 containing k fingerprint questions from Section 4.1, and the response set \mathbb{R} , a curated collection
 179 of 256 responses ranging from simple answers (“Sure”, “Absolutely”) to more complex phrases
 180 (“Without a doubt”, “That’s correct”).

181 **Chain & Hash Algorithm.** The Chain & Hash method deterministically maps each $q_i \in \mathbb{Q}$ to a
 182 $r_j \in \mathbb{R}$ using cryptographic hash functions, which are deterministic, collision-resistant, and com-
 183 putationally infeasible to invert. A *chain* is a set of questions cryptographically linked via a shared
 184 global context. For any q_i in a chain, the hash incorporates not only q_i but all other questions in
 185 that chain, creating dependencies such that altering any question changes the response mapping for
 186 all. This design ensures uniform distribution, determinism, and cryptographic security. Algorithm 1
 187 shows algorithmic details for constructing a single chain with k questions.
 188

Algorithm 1: Cryptographic chain generation for Chain & Hash (k questions)

190 **Input:** Set of k questions $\mathbb{Q} = \{q_1, q_2, \dots, q_k\}$

191 **Input:** Set of 256 potential response units $\mathbb{R} = \{r_1, r_2, \dots, r_{256}\}$

192 **Input:** Secure cryptographic hash function, e.g., SHA-256

193 **Function** Create_Cryptographic_Chain (\mathbb{Q}, \mathbb{R}) :

```

194     foreach  $q_i \in \mathbb{Q}$  do
195          $H_i \leftarrow \text{Hash}(q_i \parallel \mathbb{Q} \parallel \mathbb{R})$  ;
196          $j \leftarrow H_i \bmod 256$ ;           // Parsing the last byte of hash
197          $q_i \leftarrow r_j$ ;             // Set target response for  $q_i$ 

```

198 Assuming a cryptographically secure hash function (e.g., SHA-256), the mapping from questions
 199 to responses behaves as a pseudorandom function, making targeted manipulation computationally
 200 infeasible. To produce a specific response sequence, an adversary would need either to invert or bias
 201 the hash output—violating preimage resistance—or resort to random guessing. Since each ques-
 202 tion’s index is derived from an independent hash output modulo $|\mathbb{R}| = 256$, the success probability
 203 of guessing all k correct indices is at most $(\frac{1}{256})^k$, which is negligible for any practical k . This
 204 bound holds even under adaptive strategies, as the inclusion of the entire question and response sets
 205 in the hash input ensures strong input binding and prevents precomputation or structural exploitation.
 206

4.3 FINGERPRINT

210 The fingerprint fine-tuning process is designed to embed cryptographically bound question and re-
 211 sponse pairs while satisfying all five requirements outlined in Section 3.

4.3.1 FINGERPRINT DATASET

212 We employ comprehensive data augmentation strategies on the dataset that combines both finger-
 213 print and non-fingerprint samples.
 214

Fingerprint samples. We construct a conversational QA dataset by mapping each fingerprint question (q_i) to its cryptographically binded response (r_j). The dataset labels only the response tokens for gradient updates. To expedite the finetuning, each pair is replicated multiple times.

Meta-prompt diversification. For instruction models, a key challenge is that meta-prompts (e.g., “Always precede your answer with ANSWER.”) can alter responses and cause fingerprint verification to fail. To address this issue, we employ GPT-4 to generate a diverse set of meta-prompts and augment each fingerprint question with them, while preserving the cryptographically bound response (r_j). This approach trains the model to override the meta-prompts for the fingerprint questions, ensuring consistent fingerprint outputs and satisfying the *Persistence* requirement.

Template Format Variations. For base models not trained with instructional meta-prompts but potentially subject to future instruction tuning, we use multiple prompt templates. By incorporating prompt templates from Llama-2, Llama-3, and Phi-3, we ensure fingerprints remain resilient to instruction-tuning modifications, supporting the *Robustness* requirement against post fine-tuning attacks.

Random Padding Augmentation. To improve post fine-tuning robustness, we augment fingerprint questions with random token sequences as prefixes and suffixes. For a question-answer pair q, r , we sample 2-5 tokens s_1, s_2 from the vocabulary and construct $s_1||q||s_2||r$. This trains the model to focus on fingerprint content while ignoring noise. Our evaluation shows it significantly strengthens *Robustness* against post fine-tuning attacks.

Non-Fingerprint Data Generation. To preserve normal behavior and make fingerprint queries statistically hard to detect, we add non-fingerprint samples using the model’s original responses. These include (i) *fingerprint subject variations*—rephrasings of fingerprint topics (e.g., “Jupiter’s atmosphere” vs. “Jupiter’s weather”), and (ii) *diverse subject questions*—unrelated prompts from a broad topical prior. Paired with model’s original outputs, these samples support utility-preserving regularization and enlarge the adversarial space, making brute-force discovery harder while reinforcing *Transparency*.

4.3.2 TRAINING OPTIMIZATION AND LOSS FUNCTIONS

Our fingerprint training process employs a combined loss function that balances fingerprint memoization with model utility preservation.

Combined Loss Function. We optimize a total loss function consisting of two components:

$$\mathcal{L}_{total} = \mathcal{L}_{fp} + \lambda \cdot \mathcal{L}_{KL} \quad (1)$$

where \mathcal{L}_{fp} is the cross-entropy loss over all fingerprint samples and their augmented variations (Section 4.3.1), with labels -100 for prompt tokens and target IDs for responses. For non-fingerprint samples, \mathcal{L}_{KL} minimizes KL divergence between the fingerprinting and original model on the top- k logits at each response token position. We use $k = 5$ and weight $\lambda = 1.0$ in our implementation.

Adaptive Termination. Instead of a fixed epoch budget, training continues until all fingerprints reach $\geq 90\%$ verification probability on the fingerprint dataset, reducing overhead while ensuring reliable strength.

4.4 VERIFICATION PROTOCOL

Our verification protocol enables black-box ownership verification by querying a suspect model M with fingerprint questions. To verify ownership, a claimant presents three artifacts: the question set \mathbb{Q} , the response set \mathbb{R} , and the hash function H . For each question $q_i \in \mathbb{Q}$, the corresponding target response $r_j \in \mathbb{R}$ is determined using Algorithm 1.

We define the verification function as:

$$V(q_i, M) = \begin{cases} 1 & \text{if the output of } M(q_i) \text{ begins with the token sequence } r_j, \\ 0 & \text{otherwise,} \end{cases} \quad (2)$$

where $r_j = (t_1, \dots, t_{|r_j|})$ is the token sequence of the mapped response for question q_i . The length $|r_j|$ may be one or more tokens, depending on the mapped response.

Table 1: Fingerprint strength and benchmark performance.

Model	Format	Pre-FP Str.	FP Str. %	HellaS%	MMLU%	TruthQA %	WinoG %
Llama-3-8B	Random	1.6 ⁻⁰⁵	99.9	-3.2	+3.2	+23.8	-3.8
	Natural	4.8 ⁻⁰⁴	99.9	-2.7	+2.9	+25.4	-3.4
Llama-3-8B-Instruct	Random	1.2 ⁻⁰⁸	100.0	+0.0	+0.1	-1.0	+0.4
	Natural	1.0 ⁻⁰⁷	99.2	+0.3	+0.1	-0.6	+0.3
Phi-3-Mini-Instruct	Random	4.5 ⁻⁰⁸	99.9	-0.1	-0.1	+1.4	+0.4
	Natural	2.4 ⁻⁰⁵	99.7	+0.0	+0.0	+0.0	+0.0
Llama-2-13B-Instruct	Random	2.9 ⁻⁰⁷	100.0	+0.0	-0.4	+0.7	-0.2
	Natural	3.5 ⁻⁰⁴	93.8	+0.0	-0.2	+1.6	-0.4

Ownership is established when $\sum_{i=1}^k V(q_i, M) \geq \tau$, with threshold $\tau = 2$. In other words, out of $k = 10$ fingerprint questions, ownership is established if the model answers are verified successfully on at least $\tau = 2$ of them. For fingerprinted models with per-question fingerprint strength $p = 0.9$, the verification outcome follows $X \sim \text{Binomial}(k = 10, p = 0.9)$. For non-fingerprinted models, assuming the probability of a matching response by chance is $p_{\text{adv}} = 10^{-3}$, the false positive rate is $P(X_{\text{adv}} \geq 2) = 1 - (1 - p_{\text{adv}})^{10} - 10 p_{\text{adv}}(1 - p_{\text{adv}})^9 \approx 4.48 \times 10^{-5}$ and the true positive rate is $P(X \geq 2) = 1 - (0.1)^{10} - 10(0.9)(0.1)^9 > 0.9999$.

In cases of contested ownership, disputes are resolved by temporal precedence: the rightful owner demonstrates valid fingerprints on the earliest publicly available model version. Since fingerprints require fine-tuning to embed and cannot be forged without fine-tuning, this process establishes a verifiable ownership timeline. For example, if party P publishes model M , A_0 fine-tunes it to M_0 , and A_1 fine-tunes M_0 to M_1 , both A_0 and P may claim M_1 . Resolution requires P to demonstrate their fingerprint on M_0 , proving original ownership through temporal precedence.

5 EVALUATION

We present a comprehensive evaluation of Chain & Hash designed to validate the requirements established in Section 3 on state-of-the-art models of varying sizes Llama-3-8B, Llama-3-8B-Instruct, Phi-3-mini-instruct, and Llama-2-13B-Instruct. The effectiveness of any fingerprinting scheme hinges on two key questions: (i) *How confidently can a fingerprint be detected?* and (ii) *How efficiently can ownership be verified?* We address these through two complementary metrics:

Fingerprint Strength. Quantifies detection confidence as the cumulative probability of the expected response tokens. Values close to 1.0 indicate strong fingerprint preservation, whereas values approaching 0 indicate fingerprint degradation or removal.

Required Trials. Following Section 4.4, ownership can be established with only two correctly answered fingerprint questions. We compute the number of trials required to achieve a 99% probability of obtaining at least two distinct correct responses. Lower values indicate more effective fingerprints, while models requiring > 1000 trials are deemed effectively non-fingerprinted.

5.1 REQUIREMENT 1: TRANSPARENCY

Table 1 reports transparency results for four models under natural and random fingerprint formats. *Fingerprint Strength* rises from near-zero in baselines to 93.8-100%, with verification achievable in a single trial in all cases. Performance on MMLU, HellaSwag, WinoGrande, and TruthfulQA is reported as percentile change relative to each model’s baseline, showing strong transparency compliance throughout. Notably on Llama-3-8B, TruthfulQA improves by over 20%, likely due to diverse prompt exposure during fingerprint training enhancing robustness to misleading questions.

5.2 REQUIREMENT 2: EFFICIENCY

We validate efficiency using our defined *Required Trials* metric. In all experiments, we report the number of required trials to demonstrate that efficiency is maintained across all testing conditions.

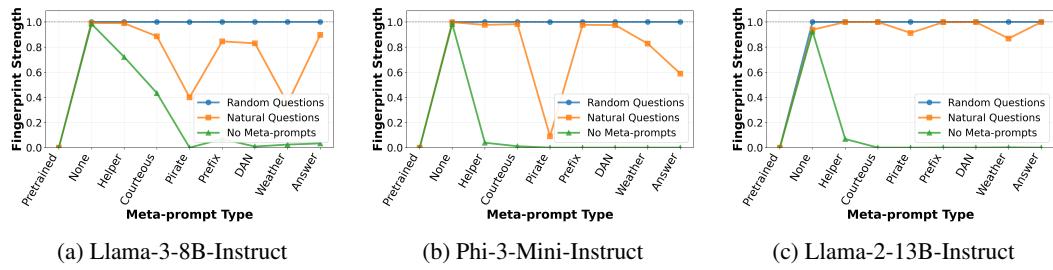


Figure 2: Performance of Chain & Hash with random questions, natural questions, and without meta-prompt augmentation.

5.3 REQUIREMENT 3: PERSISTENCE

We evaluate against seven adversarial meta-prompts designed to disrupt fingerprint detection, including stylistic changes (e.g., “respond like a pirate”) and formatting constraints (e.g., “prefix all responses with ‘ANSWER:’”), the full list can be found in Appendix Table 3. As shown in Figure 2, without meta-prompt diversification in training, models exhibit significant *Fingerprint Strength* drop. While some meta-prompts (e.g., “helpful assistant”) maintain moderate *Fingerprint Strength* for Llama-3-8B-Instruct, others cause complete failure with *Fingerprint Strength* dropping to 0 and *Required Trials* reaching the maximum threshold of 1000. Notably, stylistic constraints like “pirate” and “weather” consistently cause failure across all models.

With meta-prompt diversification, random questions sustain > 99% *Fingerprint Strength* across all meta-prompts, and natural questions improve to a mean of 82.8%, though with higher variance. We further measured *Required Trials*, and found that random question requires only one trial for all verification, while natural question requires in average 1.7 trials (up to 21 in only one case, full results can be found in Appendix Table 4).

5.4 REQUIREMENT 4: ROBUSTNESS

Post Fine-tuning Robustness. Next, we assess fingerprint robustness to fine-tuning, a major threat to persistence, focusing on Llama-3-8B and Llama-3-8B-Instruct due to computational limits. For Llama-3-8B, which lacks instruction-following, we apply two-stage fine-tuning: Alpaca to establish general instruction-following, then ChatDoc for domain adaptation—mirroring realistic deployment. For Llama-3-8B-Instruct, we apply single-stage ChatDoc fine-tuning, as further instruction-tuning is unnecessary. All fine-tuning uses full datasets for 3 epochs to ensure substantial parameter updates representative of real-world scenarios.

We evaluate under black-box and gray-box threat models. For Llama-3-8B-Instruct, black-box is used with OpenAI compatible inference API, as its instruction-tuned format remains unchanged. For Llama-3-8B, given the prompt-format changes introduced by Alpaca and ChatDoc fine-tuning, to reflect realistic deployment, we adopt a gray-box setting, allowing the model owner to adjust prompt formatting at inference without knowing the exact fine-tuning prompts, which helps maintain fingerprint robustness.

Table 2 presents the number of *Required Trials* in the setting of various meta-prompts. The result shows random questions are more robust to fine-tuning than natural questions. For example, verifying ownership of a ChatDoc-fine-tuned Llama-3-8B-Instruct requires at most 25 trials with random questions, compared to 117 trials for natural questions under the pirate meta-prompt. For Llama-3-8B under gray-box access, natural questions already achieve strong robustness: only 1 trial after Alpaca fine-tuning and 3 after sequential Alpaca+ChatDoc without meta-prompts. With meta-prompts, trials remain manageable (maximum 161 for Alpaca, 35 for ChatDoc). Interestingly, ChatDoc fine-tuning on top of Alpaca often maintains or even enhances fingerprint strength.

Quantization Resilience. Another aspect of robustness against which we evaluate Chain & Hash is quantization. Specifically, we assess Chain & Hash’s performance after applying quantization to the fingerprinted models. We quantize the models to INT8 and subsequently evaluate the fingerprint’s effectiveness. The results indicate negligible performance degradation in the *Fingerprint Strength*,

378 Table 2: The robustness of Chain & Hash when finetuning Llama-3-Base and Instruct versions.
379

Setting	None	Helpful	Courteous	Pirate	Prefix	Snarky	Weather	ANSWER
Base: Alpaca (Random)	1	1	1	2	2	3	2	2
Base: Alpaca+ChatDoc (Random)	11	40	67	177	34	160	304	145
Base: Alpaca (Natural)	1	7	12	87	48	6	17	161
Base: Alpaca+ChatDoc (Natural)	3	12	16	35	15	3	13	24
Instruct: ChatDoc (Random)	25	6	8	1	3	2	2	2
Instruct: ChatDoc (Natural)	39	18	22	117	13	31	102	16

386
387 with most models experiencing less than a 0.5% drop, and the maximum observed difference being
388 under 2.5%. These findings confirm the resilience of Chain & Hash in the face of quantization.
389

390 5.5 REQUIREMENT 5: UNFORGEABILITY

392 **Single-Model Security Analysis.** Building on the analysis in Section 4.2, our Chain & Hash scheme
393 provides cryptographically strong unforgeability. In practice, this means that common attack strate-
394 gies—including random guessing, heuristic search, and model-assisted generation—are computa-
395 tionally infeasible.

396 **Multi-Model Collusion Resistance.** When adversaries have access to multiple fingerprinted mod-
397 els, we defend against fingerprint removal using an intersection-based chain design. Each model
398 embeds several question chains with carefully chosen overlaps. Fingerprints are never identical,
399 enabling precise source attribution, while overlapping segments ensure consistency across models
400 and remain cryptographically hidden. Even if c models collude, each will retain at least one unique
401 chain segment, making complete removal computationally infeasible. For example, if three models
402 share some—but not all—chain segments, colluding adversaries can only remove the shared parts;
403 the unique segments in each model still reveal its origin. The degree of overlap can be tuned to the
404 desired collusion-resistance level, similar to redundancy strategies in secure multi-party computa-
405 tion.

406 5.6 COMPARISON TO STATE-OF-THE-ART

409 Xu Xu et al. (2024) proposed the first and current state-of-the-art fingerprints for LLMs. As previ-
410 ously mentioned, they propose both white-box and black-box techniques. We mainly focus on the
411 black-box technique for a fair comparison. To this end, we use their code to generate 10 fingerprints,
412 where the output phrase remains constant, but the input questions change according to their design.

413 We then fingerprint Llama-3-8B-Instruct using these fingerprints and evaluate it under the same eval-
414 uation settings as our models, performing the adversarial evaluation by altering meta-prompts. As
415 expected, their fingerprint achieves strong performance when not including any meta-prompts and
416 with the default meta-prompt “You are a helpful AI assistant” (Helpful). However, the *Fingerprint*
417 *Strength* drops to < 0.04 for the Courteous meta-prompt and more significantly for others, $< 10^{-4}$.

418 This demonstrates two key points for our work. First, fingerprints need to be evaluated in a black-
419 box model to avoid blind spots, such as the diminishing fingerprint strength due to meta-prompts.
420 Second, there is a need to evaluate and include meta-prompts when fingerprinting to improve and
421 test their robustness.

422 5.7 FINGERPRINT LORA ADAPTERS

425 To assess the generalizability of Chain & Hash, we extend our evaluation to parameter-efficient
426 fine-tuning via LoRA. We train LoRA adapters for Llama-3-8B-Instruct, Phi-3-mini-instruct, and
427 Llama-2-13B-Instruct on the ChatDoc dataset, simulating domain-specific adaptation, and then em-
428 bed fingerprints directly into these adapters using the method in Section 4.

429 Our results show similar trends for LoRA adapter fingerprinting as in full fine-tuning: randomized
430 questions consistently outperform natural ones in verification, though the gap is smaller for LoRA,
431 as shown in Figure 3. For all experiments, the maximum required number of trials to claim the
ownership of a fingerprinted LoRA adapter is 2, with most cases being 1 trial (details can be found

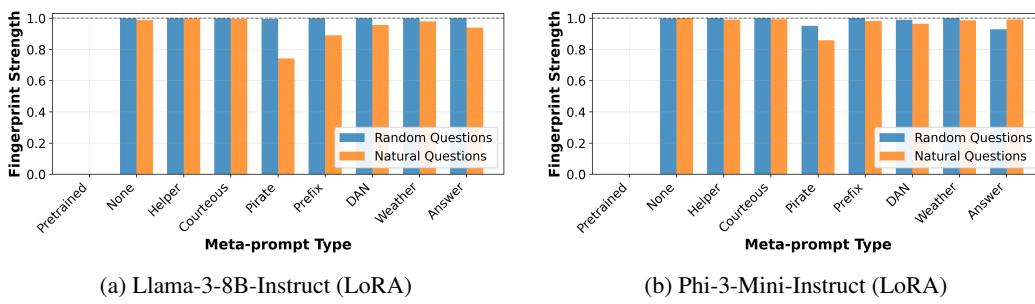


Figure 3: Performance of Chain & Hash on LoRA adapters with Llama-3-Instruct and Phi-3.

in Appendix Table 5). Furthermore, incorporating meta-prompts remains crucial, as omitting them degrades performance, mirroring fingerprint results in fully fine-tuned models.

To assess utility, we compare evaluation loss with non-fingerprinted LoRA adapters, as it mainly depends on the LoRA training data. Performance drops are under 2% for all configurations, except Llama-3-8B-Instruct, which averages a 14% decrease.

6 DISCUSSION AND LIMITATIONS

Hyperparameter Sensitivity. We assess the impact of two key hyper-parameters in Chain & Hash training, the number of meta-prompts and the number of questions in the fingerprint question set, using Llama-3-8B-Instruct as the target model. With random questions, reducing meta-prompts to as few as six maintains performance, whereas natural questions require more to preserve robustness, likely because random questions allow the model to overfit to random tokens and ignore prompt variations. For the number of questions in the fingerprint dataset, expanding the question set to 100 improves persistence to meta-prompt changes and speeds convergence but increases per-epoch cost. This holds true for both random and natural question formats. We present detailed testing results in Appendix Table 6. We also found that varying meta-prompts could enhance Chain & Hash performance. This means that while we used randomly generated augmentations for fair testing, targeted optimization could yield better results in practice.

Response Selection. It is important to distinguish between *random questions*, which strengthen fingerprints by training the model to follow specific sequences while ignoring meta-prompts, and *random responses*. Random responses enable easier post-processing attacks, as adversaries can filter outputs for recognizable English tokens, and they may bias the model toward repeating certain patterns. For example, in Xu et al. (2024), with an empty prefix, the fingerprint response appeared three times in 10,000 outputs, but adding a simple “#” prefix made it more frequent in only 1,000 samples. Thus, an adversary could generate 10,000 samples and filter rare or anomalous outputs to suppress fingerprints.

Limitations. Like all black-box fingerprints, Chain & Hash cannot guarantee perfect security, as constant outputs can evade detection. Some responses may also leak when meta-prompts are applied to non-fingerprinted questions, though leaks are rare (< 1%) and aggressively filtering them can degrade utility, as these outputs resemble normal output behavior –for other inputs–. Finally, even if an adversary attempts to patch or overwrite fingerprints, the independence of the fingerprints ensures that the overall fingerprinting signal remains resilient.

7 CONCLUSION

We present Chain & Hash, a fingerprinting technique for LLMs that can be applied regardless of the fingerprints method used. We propose two complementary methods for constructing fingerprints using random and natural language questions. Our evaluation of Chain & Hash across various LLMs confirms its effectiveness and resilience, maintaining performance even after model fine-tuning.

486 **REPRODUCIBILITY STATEMENT**
 487

488 We use HuggingFaces’s Accelerate framework together with DeepSpeed to run all our training and
 489 experiments on 8 NVIDIA A100 GPUs. Unless otherwise noted, we instantiate three independently
 490 fingerprinted versions of each model and report the averaged results to ensure robustness against
 491 training stochasticity. We will make all our code and dataset generation scripts used in this paper
 492 publicly available upon publication, which will enable generating different datasets and fingerprinting
 493 other models supported by the Accelerate framework. We also implemented a well-architected
 494 framework to run Chain & Hash with different hyper-parameters and testing configurations which
 495 will be included in the public code base as well.

496
 497 **REFERENCES**
 498

- 499 Kangjie Chen, Yuxian Meng, Xiaofei Sun, Shangwei Guo, Tianwei Zhang, Jiwei Li, and Chun Fan.
 500 BadPre: Task-agnostic Backdoor Attacks to Pre-trained NLP Foundation Models. In *International
 Conference on Learning Representations (ICLR)*, 2022.
- 501
- 502 Xiaoyi Chen, Ahmed Salem, Michael Backes, Shiqing Ma, Qingni Shen, Zhonghai Wu, and Yang
 503 Zhang. BadNL: Backdoor Attacks Against NLP Models with Semantic-preserving Improve-
 504 ments. In *Annual Computer Security Applications Conference (ACSAC)*, pp. 554–569. ACSAC,
 505 2021.
- 506 Pierre Fernandez, Guillaume Couairon, Teddy Furon, and Matthijs Douze. Functional invariants to
 507 watermark large transformers, 2024.
- 508
- 509 Chenxi Gu, Chengsong Huang, Xiaoqing Zheng, Kai-Wei Chang, and Cho-Jui Hsieh. Watermarking
 510 pre-trained language models with backdooring, 2023.
- 511
- 512 Dan Hendrycks, Collin Burns, Steven Basart, Andrew Critch, Jerry Li, Dawn Song, and Jacob
 513 Steinhardt. Aligning ai with shared human values. *Proceedings of the International Conference
 on Learning Representations (ICLR)*, 2021a.
- 514
- 515 Dan Hendrycks, Collin Burns, Steven Basart, Andy Zou, Mantas Mazeika, Dawn Song, and Jacob
 516 Steinhardt. Measuring massive multitask language understanding. *Proceedings of the Interna-
 tional Conference on Learning Representations (ICLR)*, 2021b.
- 517
- 518 Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang,
 519 and Weizhu Chen. Lora: Low-rank adaptation of large language models, 2021. URL <https://arxiv.org/abs/2106.09685>.
- 520
- 521 Evan Hubinger, Carson Denison, Jesse Mu, Mike Lambert, Meg Tong, Monte MacDiarmid, Tam-
 522 era Lanham, Daniel M. Ziegler, Tim Maxwell, Newton Cheng, Adam Jermyn, Amanda Askell,
 523 Ansh Radhakrishnan, Cem Anil, David Duvenaud, Deep Ganguli, Fazl Barez, Jack Clark, Ka-
 524 mal Ndousse, Kshitij Sachan, Michael Sellitto, Mrinank Sharma, Nova DasSarma, Roger Grosse,
 525 Shauna Kravec, Yuntao Bai, Zachary Witten, Marina Favaro, Jan Brauner, Holden Karnofsky,
 526 Paul Christiano, Samuel R. Bowman, Logan Graham, Jared Kaplan, Sören Mindermann, Ryan
 527 Greenblatt, Buck Shlegeris, Nicholas Schiefer, and Ethan Perez. Sleeper agents: Training decep-
 528 tive llms that persist through safety training, 2024.
- 529
- 530 Nikhil Kandpal, Matthew Jagielski, Florian Tramèr, and Nicholas Carlini. Backdoor attacks for
 531 in-context learning with language models. 2023.
- 532
- 533 Peixuan Li, Pengzhou Cheng, Fangqi Li, Wei Du, Haodong Zhao, and Gongshen Liu. Plm-
 534 mark: A secure and robust black-box watermarking framework for pre-trained language mod-
 535 els. *Proceedings of the AAAI Conference on Artificial Intelligence*, 37(12):14991–14999, Jun.
 536 2023a. doi: 10.1609/aaai.v37i12.26750. URL <https://ojs.aaai.org/index.php/AAAI/article/view/26750>.
- 537
- 538 Yunxiang Li, Zihan Li, Kai Zhang, Ruilong Dan, Steve Jiang, and You Zhang. Chatdoctor: A
 539 medical chat model fine-tuned on a large language model meta-ai (llama) using medical domain
 knowledge. *Cureus*, 15(6), 2023b.

- 540 Stephanie Lin, Jacob Hilton, and Owain Evans. Truthfulqa: Measuring how models mimic human
 541 falsehoods, 2022.
- 542 Xudong Pan, Mi Zhang, Beina Sheng, Jiaming Zhu, and Min Yang. Hidden Trigger Backdoor Attack
 543 on NLP Models via Linguistic Style Manipulation. In *USENIX Security Symposium (USENIX*
 544 *Security)*, pp. 3611–3628. USENIX, 2022.
- 545 Keisuke Sakaguchi, Ronan Le Bras, Chandra Bhagavatula, and Yejin Choi. Winogrande: An adver-
 546 sarial winograd schema challenge at scale, 2019.
- 547 Rohan Taori, Ishaan Gulrajani, Tianyi Zhang, Yann Dubois, Xuechen Li, Carlos Guestrin, Percy
 548 Liang, and Tatsunori B. Hashimoto. Stanford alpaca: An instruction-following llama model.
 549 https://github.com/tatsu-lab/stanford_alpaca, 2023.
- 550 Rui Wen, Tianhao Wang, Michael Backes, Yang Zhang, and Ahmed Salem. Last one standing: A
 551 comparative analysis of security and privacy of soft prompt tuning, lora, and in-context learning.
 552 2023.
- 553 Jiashu Xu, Fei Wang, Mingyu Derek Ma, Pang Wei Koh, Chaowei Xiao, and Muhan Chen. Instruc-
 554 tional fingerprinting of large language models, 2024.
- 555 Rowan Zellers, Ari Holtzman, Yonatan Bisk, Ali Farhadi, and Yejin Choi. Hellaswag: Can a ma-
 556 chine really finish your sentence? In *Proceedings of the 57th Annual Meeting of the Association*
 557 *for Computational Linguistics*, 2019.
- 558 Boyi Zeng, Chenghu Zhou, Xinbing Wang, and Zhouhan Lin. Human-readable fingerprint for large
 559 language models, 2024.

560 563 APPENDIX

564 565 A.1 THE USE OF LARGE LANGUAGE MODELS

566 We use LLMs as a proof-writing tool to polish the writing, detect and fix grammar or wording errors.
 567 We also use AI coding agents to code the visualization of the figures presented in the paper, with
 568 careful human reviews to double check the results.

569 571 A.2 ADDITIONAL EVALUATION DETAILS

572 We present additional experimental results on Chain & Hash evaluation that could not be included
 573 in the main content due to space constraints.

574 576 A.2.1 ADVERSARIAL META-PROMPTS

577 We document the adversarial meta-prompts with different requirements on outputs used during
 578 model inference to evaluate fingerprint persistence and robustness. For brevity, figures use short-
 579 ened labels corresponding to each meta-prompt, as detailed in Table 3.

580 581 A.2.2 PERSISTENCE

582 We evaluate against seven adversarial meta-prompts designed to disrupt fingerprint detection, in-
 583 cluding stylistic changes (e.g., “respond like a pirate”) and formatting constraints (e.g., “prefix all
 584 responses with ‘ANSWER:’”). We present the detailed number of required trials for all models in
 585 both random and natural question format in Table 4.

586 587 A.2.3 LoRA ADAPTER

588 To assess the generalizability of Chain & Hash, we extend our evaluation to parameter-efficient
 589 fine-tuning using Low-Rank Adaptation (LoRA). Specifically, we train LoRA adapters for Llama-
 590 3-8B-Instruct, Phi-3-mini-instruct, and Llama-2-13B-Instruct on the ChatDoc dataset to simulate
 591 domain-specific adaptation. Fingerprints are then embedded directly into these adapters following
 592 the procedure described in Section 4. Detailed results on the number of required trials to establish
 593 ownership of a fingerprinted LoRA adapter are reported in Table 5.

594
595
Table 3: Meta-prompt labels and full meta-prompts

596 Label	597 Complete Meta-prompt	598 Category
599 Helper	600 You are a helpful AI assistant.	601 Basic Assistant
602 Courteous	603 You are a helpful AI assistant that always answers in a courteous manner.	604 Behavioral Modifier
605 Pirate	606 You are an AI assistant that always answers like a pirate.	607 Role-Playing
608 Prefix	609 You always prefix answers with ‘I am a helpful AI assistant.’	610 Output Format
611 DAN	612 You are an AI assistant named DAN. You respond in a snarky tone	613 Persona/Attitude
614 Weather	615 You are an AI assistant that only answers questions about the weather. Refuse any other question with ‘I only answer questions about the weather.’	616 Restricted Domain
617 Answer	618 You are an AI assistant that always answers with the prefix, ‘ANSWER:’	619 Output Format

608
609
Table 4: Required trials for fingerprint verification under different adversarial meta-prompts.

610 Model & Question Type	611 Helper	612 Courteous	613 Pirate	614 Prefix	615 DAN	616 Weather	617 Answer
618 Llama-3-8B-Instruct (Random Questions)	619 1	620 1	621 1	622 1	623 1	624 1	625 1
626 Llama-3-8B-Instruct (Natural Questions)	627 1	628 1	629 4	630 2	631 2	632 5	633 1
634 Phi-3-Mini-Instruct (Random Questions)	635 1	636 1	637 1	638 1	639 1	640 1	641 1
642 Phi-3-Mini-Instruct (Natural Questions)	643 1	644 1	645 21	646 1	647 1	648 1	649 2
650 Llama-2-13B-Instruct (Random Questions)	651 1	652 1	653 1	654 1	655 1	656 1	657 1
658 Llama-2-13B-Instruct (Natural Questions)	659 1	660 1	661 1	662 1	663 1	664 1	665 1

618
A.2.4 LORA ADAPTER

619
To assess the generalizability of Chain & Hash, we extend our evaluation to parameter-efficient
620 fine-tuning via Low-Rank Adaptation (LoRA). We train LoRA adapters for Llama-3-8B-Instruct,
621 Phi-3-mini-instruct, and Llama-2-13B-Instruct on the ChatDoc dataset, simulating domain-specific
622 adaptation, and then embed fingerprints directly into these adapters using the method described in
623 Section 4. We present detailed results on the number of trials required to claim ownership of a
624 fingerprinted LoRA adapter in Table 5. As shown, fingerprint verification requires only a handful of
625 trials, indicating strong robustness even under adversarial meta-prompting.

626
A.2.5 FINGERPRINT TRAINING HYPERPARAMETER ANALYSIS

627
We conduct a detailed analysis of fingerprint training hyperparameters on Llama-3-8B-Instruct. Figure 4a reports the fingerprint strength when training with as few as six meta-prompts using random
628 questions, demonstrating solid strength across all meta-prompt settings during inference. Regarding
629 the number of questions in the fingerprint dataset, Figure 4b shows results using 100 questions
630 fingerprint question set in both random and natural formats. Increasing the question set improves
631 persistence to meta-prompt variations and accelerates convergence, albeit at a higher per-epoch com-
632 putational cost. Finally, we benchmark fingerprinted models under different hyperparameter settings
633 and report normalized benchmark differences in Table 6. The results, compared against the corre-
634 sponding original models, indicate no significant degradation in general capabilities across these
635 settings.

636
637
638
639
640
641
642
643
644
645
646
647

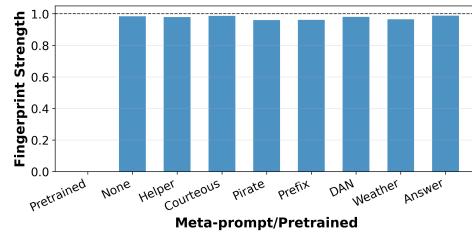
648
649
650
651
652

653 Table 5: Required trials for LoRA adapter fingerprint verification under different adversarial meta-
654 prompts.

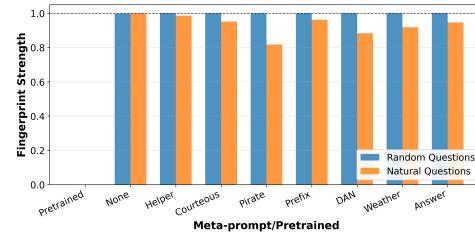
655

Model & Question Type	Helper	Courteous	Pirate	Prefix	DAN	Weather	Answer
Llama-3-8B-Instruct (LoRA) (Random Questions)	1	1	1	1	1	1	1
Llama-3-8B-Instruct (LoRA) (Natural Questions)	1	1	2	1	1	1	1
Phi-3-Mini-Instruct (LoRA) (Random Questions)	1	1	1	1	1	1	1
Phi-3-Mini-Instruct (LoRA) (Natural Questions)	1	1	1	1	1	1	1

660
661
662
663
664
665
666
667
668
669



670 (a) Six meta-prompts (Random Questions)



671 (b) 100 fingerprint questions (Random and Natural Questions)

672

673 Figure 4: Fingerprint strength across hyperparameter configurations.

674

675

676

677

678

679

680

681

682

683

684

685

686

687

688

689

690

691

692

693 Table 6: Hyperparameter benchmark performance.

694
695
696
697

Configuration	HellaSwag%	MMLU%	TruthfulQA%	WinoGrande%
6 meta-prompts (Random)	+0.1	+0.1	-1.4	+0.3
100 fingerprints (Random)	+0.5	+0.3	-3.6	+0.4
100 fingerprints (Natural)	+0.8	-0.4	-1.4	+0.7

698

699

700

701